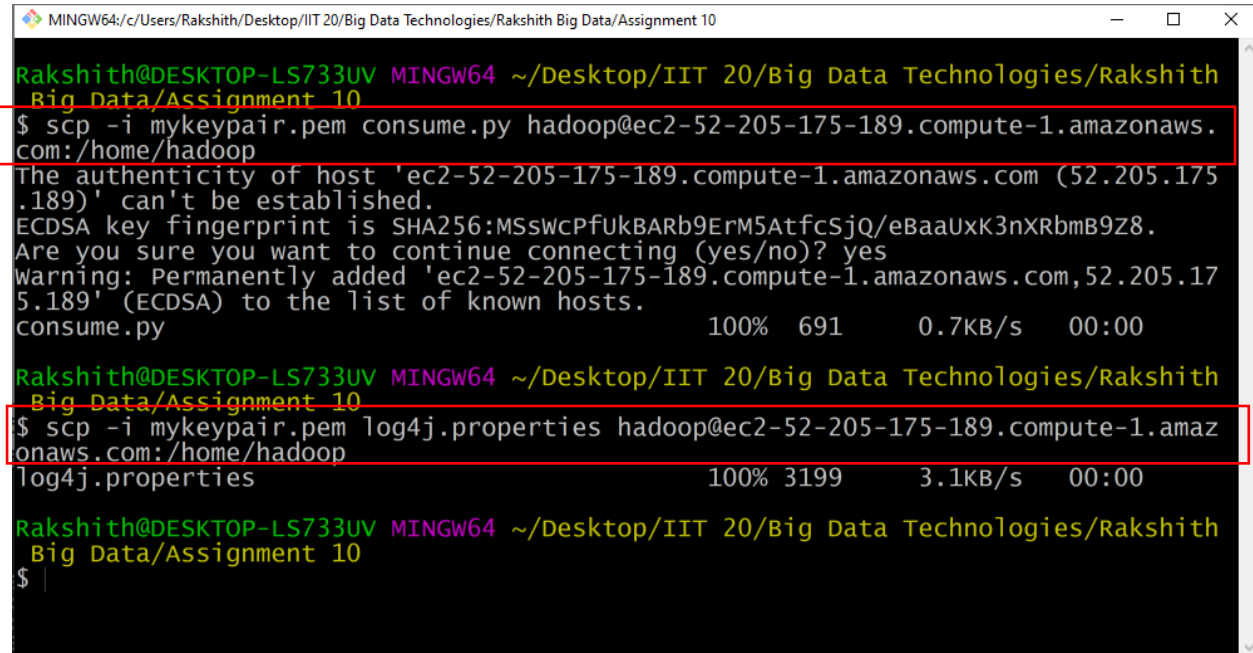


SCP Command:



```

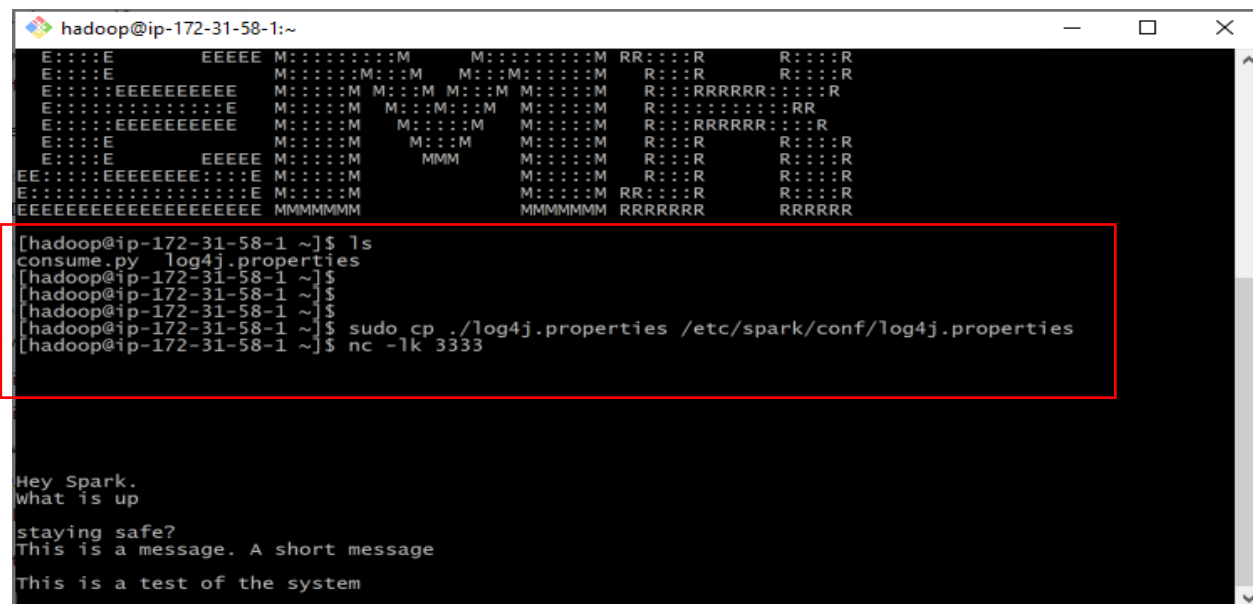
MINGW64:/c:/Users/Rakshith/Desktop/IIT 20/Big Data Technologies/Rakshith Big Data/Assignment 10
Rakshith@DESKTOP-LS733UV MINGW64 ~/Desktop/IIT 20/Big Data Technologies/Rakshith Big Data/Assignment 10
$ scp -i mykeypair.pem consume.py hadoop@ec2-52-205-175-189.compute-1.amazonaws.com:/home/hadoop
The authenticity of host 'ec2-52-205-175-189.compute-1.amazonaws.com (52.205.175.189)' can't be established.
ECDSA key fingerprint is SHA256:MSSwCPfUkBARb9ErM5AtfcSjQ/eBaaUxK3nXRbmB9Z8.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'ec2-52-205-175-189.compute-1.amazonaws.com,52.205.175.189' (ECDSA) to the list of known hosts.
consume.py                                100% 691      0.7KB/s   00:00

Rakshith@DESKTOP-LS733UV MINGW64 ~/Desktop/IIT 20/Big Data Technologies/Rakshith Big Data/Assignment 10
$ scp -i mykeypair.pem log4j.properties hadoop@ec2-52-205-175-189.compute-1.amazonaws.com:/home/hadoop
log4j.properties                          100% 3199      3.1KB/s   00:00

Rakshith@DESKTOP-LS733UV MINGW64 ~/Desktop/IIT 20/Big Data Technologies/Rakshith Big Data/Assignment 10
$

```

EC2-1 Window:



```

hadoop@ip-172-31-58-1:~
E:::E EEEEE M:::M M:::M RR:::R R:::R
E:::E M:::M M:::M M:::M R:::R R:::R
E:::E EEEEEEE M:::M M:::M M:::M R:::RRRRR:::R
E:::E EEEEEEE M:::M M:::M M:::M R:::RRRRR:::R
E:::E EEEEEEE M:::M M:::M M:::M R:::RRRRR:::R
E:::E EEEEE M:::M M:::M M:::M R:::R R:::R
E:::E EEEEE M:::M M:::M M:::M R:::R R:::R
EE:::EEEEEE:::E M:::M M:::M M:::M R:::R R:::R
E:::E EEEEEEE M:::M M:::M RR:::R R:::R
EEEEEEEEEEEEEEEE M:::M M:::M RR:::R R:::R

[hadoop@ip-172-31-58-1 ~]$ ls
consume.py log4j.properties
[hadoop@ip-172-31-58-1 ~]$
[hadoop@ip-172-31-58-1 ~]$
[hadoop@ip-172-31-58-1 ~]$ sudo cp ./log4j.properties /etc/spark/conf/log4j.properties
[hadoop@ip-172-31-58-1 ~]$ nc -lk 3333

Hey Spark.
what is up

staying safe?
This is a message. A short message

This is a test of the system

```

A20424771

EC2-2 Window:

```
hadoop@ip-172-31-58-1:~$ ls
consume.py  log4j.properties
hadoop@ip-172-31-58-1:~$ spark-submit consume.py
20/04/02 01:22:13 WARN StreamingContext: dynamicAllocation is enabled for this application. enabling dynamic allocation for spark streaming ap
plications can cause data loss if write Ahead Log is not enabled for non-replayable sources like Flume. See the programming guide for details o
n how to enable the write Ahead Log.
-----
Time: 2020-04-02 01:22:30
-----
(u'', 1)
(u'what', 1)
(u'is', 1)
(u'spark', 1)
(u'up', 1)
(u'hey', 1)
-----
Time: 2020-04-02 01:22:40
-----
Time: 2020-04-02 01:22:50
-----
(u'staying', 1)
(u'safe?', 1)
-----
Time: 2020-04-02 01:23:00
-----
Time: 2020-04-02 01:23:10
-----
Time: 2020-04-02 01:23:20
-----
Time: 2020-04-02 01:23:30
-----
(u'a', 1)
(u'A', 1)
(u'is', 1)
(u'This', 1)
(u'message', 1)
(u'message', 1)
(u'short', 1)
-----
Time: 2020-04-02 01:23:40
-----
Time: 2020-04-02 01:23:50
-----
(u'', 1)
-----
Time: 2020-04-02 01:24:00
-----
(u'a', 1)
(u'This', 1)
(u'is', 1)
(u'test', 1)
(u'the', 1)
(u'of', 1)
(u'system', 1)
-----
Time: 2020-04-02 01:24:10
-----
Time: 2020-04-02 01:24:20
-----
Time: 2020-04-02 01:24:30
-----
Time: 2020-04-02 01:24:40
```

A20424771

All in one:

The screenshot displays the AWS EMR console interface for a cluster named "My cluster 10". The cluster is in a "Waiting" state, indicating it is ready after the last step completed. The console shows various tabs for cluster management, including Summary, Application history, Monitoring, Hardware, and Configurations. The Master public DNS is listed as ec2-52-205-175-189.compute-1.amazonaws.com, and the SSH connection is available.

Overlaid on the console are two terminal windows. The left terminal window shows the output of a series of commands executed on a node in the cluster:

```
hadoop@ip-172-31-58-1 ~$ ls
consume.py  log4j.properties
hadoop@ip-172-31-58-1 ~$ cd /etc/spark/conf/
hadoop@ip-172-31-58-1 ~$ cp ./log4j.properties /etc/spark/conf/log4j.properties
hadoop@ip-172-31-58-1 ~$ nc -lk 3333
```

The right terminal window displays the logs of the cluster, showing a series of messages from the Spark application, including "Hey Spark.", "What is up", "staying safe?", "This is a message. A short message", and "This is a test of the system". The logs also show the time of each message, ranging from 2020-04-02 01:22:30 to 2020-04-02 01:24:20.